

Organization Capabilities Report

Date: 2026-05-30

Organizations: alex-place, human-flourishing-frameworks

Scope: Theoretical and novel capabilities across entire organizations

Status: Comprehensive capability mapping and integration analysis

Executive Summary

This report documents the theoretical and novel capabilities across the alex-place and human-flourishing-frameworks GitHub organizations. The capabilities span local-first AI control planes, agent orchestration, game development, business services, and human flourishing frameworks.

Total Organizations: 2

Total Repositories: 43 (alex-place) + 1 (human-flourishing-frameworks) = 44

Primary Control Plane: Lantern OS (alex-place/lantern-os)

Primary Framework: COMET LEAP (human-flourishing-frameworks/human-flourishing-frameworks)

Organization Overview

alex-place Organization

Mission: Build local-first AI control planes, game development tools, and business services with strict safety boundaries and human trial readiness gates.

- Core Capabilities:**
- Local-first AI control plane (Lantern OS)
 - Agent orchestration and MCP connectors
 - Game development (GameMaker, libGDX)
 - Business services and payment processing
 - Symbolic sandbox and RAG systems

human-flourishing-frameworks Organization

Mission: Develop frameworks for human flourishing, COMET LEAP methodology, and evidence-based decision making.

- Core Capabilities:**
- COMET LEAP framework
 - Human flourishing metrics
 - Evidence-based decision systems
 - Bayesian world modeling

- Founder wisdom protocols

Theoretical Capabilities

Local-First AI Control Plane (Lantern OS)

Theoretical Capability: Autonomous local AI control with human-in-the-loop safety gates

Components:

- ASI Arc Reactor MK1 confidence model with Brier-style error tracking
- Windsurf hooks for command validation and audit logging
- MCP canary validation for safe tool exposure
- Human trial demo readiness gates (evidence, safety, consent, rollback)
- Distributed fleet governance with agent network supervision

Theoretical Applications:

- Personal AI assistant with local privacy
- Family coordination and decision support
- Small business automation with safety boundaries
- Educational tools with capability honesty
- Creative work augmentation without replacement

Novel Aspects:

- Strict claim boundaries (no unverified ASI capability claims)
- Architecture references only (ASI patterns as design guidance, not capability)
- Brier-style error tracking for confidence calibration
- Automated safety gates with human approval requirements
- Evidence-based promotion from theoretical to proven

Agent Orchestration (gm-agent-orchestrator)

Theoretical Capability: Multi-agent coordination with MCP connectors and safety validation

Components:

- MCP (Model Context Protocol) connectors
- Agent fleet management
- Task pooling and batching
- Safety hooks and validation
- Audit logging for all agent interactions

Theoretical Applications:

- Automated research and synthesis

- Multi-step task execution with human approval
- Distributed problem solving
- Parallel processing with coordination
- Fallback and recovery paths

****Novel Aspects:****

- MCP canary validation before tool exposure
- Pre/post hooks for all agent actions
- Audit trail for transparency
- Break-glass recovery paths
- Human approval gates for sensitive operations

Game Development Tools (GameMaker, libGDX)

****Theoretical Capability:**** Game development with cloud-access enabled for users

****Components:****

- GameMaker room editor tooling
- libGDX framework fork
- Game projects (ChildOfLevistus, BadSloth, etc.)
- Cloud-access for user collaboration
- Asset management and preservation

****Theoretical Applications:****

- Educational game development
- Collaborative game creation
- Game development training
- Asset preservation and licensing
- Community game projects

****Novel Aspects:****


- Cloud-access enabled for users (not just developers)
- Asset preservation with rights review
- Game lane candidate classification
- Source-summary preservation
- Delta analysis for vendor forks

Business Services (moneybags, payment bridge)

****Theoretical Capability:**** Payment processing and business automation with transparency

****Components:****

- Payment bridge with Stripe integration



- Wallet/store functionality
- Service automation (US Bank integration)
- Cost transparency server
- Invoice generation and tracking

****Theoretical Applications:****

- Small business payment processing
- Service automation with cost tracking
- Invoice management
- Financial transparency
- Cash flow optimization

****Novel Aspects:****

- Cost transparency server for visibility
- US Bank integration for statement parsing
- Service automation with safety gates
- Invoice conversion and tracking
- Secrets/env review before promotion

Symbolic Sandbox (lantern-symbolic-sandbox)

****Theoretical Capability:**** Symbolic reasoning and RAG quarantine testing

****Components:****

- Symbolic reasoning engine
- RAG quarantine system
- Evidence validation
- Rights review gates
- Content classification

****Theoretical Applications:****

- Safe content intake and classification
- Rights validation before use
- Evidence-based content promotion
- Quarantine for unverified content
- Symbolic reasoning for decision support

****Novel Aspects:****

- Quarantine system for unverified content
- Rights review gates
- Evidence-based promotion
- Symbolic reasoning integration
- Content classification with metadata



Novel Capabilities

ASI Arc Reactor MK1

****Novel Capability:**** ASI architectural pattern integration without capability claims

****Innovation:****

- ASI patterns as architecture references only
- Strict claim boundaries (no local ASI capability)
- Brier-style error tracking for confidence
- Human trial readiness gates
- Evidence-based confidence calibration

****Theoretical Foundation:****

- ASI Alliance Vision Paper (coordination/governance)
- CUDOS Edition (decentralized compute)
- AAAI design whitepaper (agent networks)
- ASI Alliance papers (token governance)

****Novel Aspects:****

- Architecture vs capability distinction
- Blocked claims list (investment advice, free compute, etc.)
- Human trial demo readiness (18% baseline)
- Distributed fleet metrics (22% confidence)
- Safety automation gates

COMET LEAP Framework

****Novel Capability:**** 30-day model for founder wisdom and decision making

****Innovation:****

- Bayesian founder wisdom protocol
- True social internet strategy
- Evidence-based decision loops
- Founder decision cards
- Brier-style error tracking

****Theoretical Foundation:****

- Bayesian world modeling
- Evidence-based decision making
- Founder wisdom extraction



- 30-day sprint methodology
- Confidence calibration

****Novel Aspects:****

- 30-day model for rapid iteration
- Founder wisdom as structured data
- Bayesian confidence updating
- Evidence-based promotion
- True social internet strategy

Windsurf Hooks System

****Novel Capability:**** Pre/post hooks for all AI agent actions with safety validation

****Innovation:****

- Pre-run command validation (demo safety)
- Pre-MCP tool use validation (dangerous tools)
- Pre-write code validation (system paths)
- Pre-read code validation (sensitive files)
- Pre-user prompt validation (unsafe patterns)
- Post-action audit logging

****Theoretical Foundation:****

- Safety by design
- Audit trails for transparency
- Human approval gates
- Break-glass recovery
- Evidence-based validation

****Novel Aspects:****

- Comprehensive hook coverage
- PowerShell validation scripts
- Audit logging to central log
- Exit code 2 for blocking
- Human trial readiness integration

Convergence Loop Method

****Novel Capability:**** 12-step convergence loop for release decisions

****Innovation:****

- 12-step validation process
- Actionable vs held issues



- Promotion states
- Fix window (4 issues)
- Evidence-based promotion

****Theoretical Foundation:****

- Systems thinking
- Evidence-based decision making
- Risk mitigation
- Rollback paths
- Operator approval

****Novel Aspects:****

- 12-step process
- Actionable vs held distinction
- Fix window for prioritization
- Evidence receipts
- Convergence to master

Human Flourishing Frameworks

****Novel Capability:**** Frameworks for measuring and enabling human flourishing

****Innovation:****

- Human flourishing metrics
- Evidence-based interventions
- Founder wisdom protocols
- Bayesian world modeling
- COMET LEAP methodology

****Theoretical Foundation:****

- Positive psychology
- Evidence-based interventions
- Bayesian decision theory
- Founder wisdom extraction
- 30-day sprint methodology

****Novel Aspects:****

- Flourishing metrics
- Evidence-based promotion
- Founder wisdom as data
- Bayesian confidence
- COMET LEAP integration



Integration Capabilities

Lantern OS + COMET LEAP Integration

****Integration Point:**** ASI Arc Reactor MK1 uses COMET LEAP framework

****Capabilities:****

- Brier-style error tracking (from COMET LEAP)
- Founder wisdom protocols (from COMET LEAP)
- 30-day sprint methodology (from COMET LEAP)
- Human trial readiness gates (from Lantern OS)
- Safety automation (from Lantern OS)

****Novel Integration:****

- COMET LEAP provides the decision framework
- Lantern OS provides the safety automation
- Combined: Evidence-based AI control with safety gates

Agent Orchestration + Symbolic Sandbox

****Integration Point:**** MCP connectors use symbolic sandbox for content validation

****Capabilities:****

- MCP connectors provide tool access
- Symbolic sandbox provides content validation
- Combined: Safe tool use with content validation
- Audit logging for transparency
- Human approval gates

****Novel Integration:****

- Symbolic sandbox validates content before agent use
- MCP connectors provide safe tool exposure
- Combined: Safe agent operations with content validation

Business Services + Payment Bridge

****Integration Point:**** Business services use payment bridge for transactions

****Capabilities:****

- Business services provide automation
- Payment bridge provides transaction processing
- Cost transparency server provides visibility
- Combined: Automated business operations with financial transparency



****Novel Integration:****

- Cost transparency for all operations
- Invoice generation and tracking
- Service automation with financial visibility
- Combined: Business automation with financial control

Theoretical Applications

Personal AI Assistant

****Theoretical Application:**** Local AI assistant with privacy and safety

****Capabilities:****

- Local-first processing (no cloud dependency)
- Safety hooks for dangerous operations
- Human approval for sensitive tasks
- Audit logging for transparency
- Evidence-based confidence

****Novel Aspects:****

- Privacy by design (local-first)
- Safety by design (hooks)
- Transparency by design (audit)
- Evidence-based confidence
- Human-in-the-loop control

Family Coordination System

****Theoretical Application:**** Family decision support and coordination

****Capabilities:****

- Family task coordination
- Decision support with evidence
- Budget tracking and transparency
- Communication logging
- Conflict resolution support

****Novel Aspects:****

- Family-specific workflows
- Evidence-based decisions
- Financial transparency



- Communication audit
- Conflict resolution protocols

Small Business Automation

Theoretical Application: Small business automation with safety boundaries

Capabilities:

- Payment processing
- Service automation
- Cost tracking
- Invoice generation
- Safety gates for financial operations

Novel Aspects:

- Cost transparency
- Safety gates for financial ops
- Invoice automation
- Service automation with validation
- Financial audit trails

Educational Tools

Theoretical Application: Educational tools with capability honesty

Capabilities:

- Learning content delivery
- Progress tracking
- Capability honesty (no overclaiming)
- Evidence-based assessment
- Safety for student use

Novel Aspects:

- Capability honesty
- Evidence-based assessment
- Safety for students
- Progress tracking
- No overclaiming

Creative Work Augmentation

Theoretical Application: Creative work augmentation without replacement

Capabilities:



- Creative tool integration
- AI assistance for creative tasks
- Human-in-the-loop control
- Attribution and credit
- Safety for creative workflows

****Novel Aspects:****

- Augmentation not replacement
- Human-in-the-loop
- Attribution systems
- Creative safety
- No replacement claims

Novel Research Directions

ASI Pattern Integration Without Capability Claims

****Research Direction:**** How to integrate ASI architectural patterns without making unverified capability claims

****Novel Approach:****

- Architecture references only
- Strict claim boundaries
- Evidence-based promotion
- Brier-style error tracking
- Human trial readiness gates

****Theoretical Impact:****

- Safe ASI pattern adoption
- Evidence-based confidence
- No overclaiming
- Human trial validation
- Regulatory compliance preparation

Bayesian Founder Wisdom

****Research Direction:**** How to extract and use founder wisdom in Bayesian decision systems

****Novel Approach:****

- Founder wisdom as structured data
- Bayesian confidence updating
- Evidence-based promotion



- 30-day sprint methodology
- Decision card extraction

****Theoretical Impact:****

- Founder wisdom preservation
- Evidence-based decisions
- Rapid iteration
- Confidence calibration
- Wisdom democratization

Safety-First Agent Orchestration

****Research Direction:**** How to orchestrate agents with comprehensive safety validation

****Novel Approach:****

- Pre/post hooks for all actions
- MCP canary validation
- Audit logging
- Human approval gates
- Break-glass recovery

****Theoretical Impact:****

- Safe agent operations
- Transparency
- Human control
- Recovery paths
- Regulatory compliance

Evidence-Based Convergence

****Research Direction:**** How to converge systems using evidence-based promotion

****Novel Approach:****

- 12-step convergence loop
- Actionable vs held issues
- Fix window
- Evidence receipts
- Promotion states

****Theoretical Impact:****

- Evidence-based decisions
- Risk mitigation
- Rollback paths



- Operator approval
- System convergence

Human Flourishing Metrics

****Research Direction:**** How to measure and enable human flourishing

****Novel Approach:****

- Flourishing metrics
- Evidence-based interventions
- Founder wisdom protocols
- Bayesian world modeling
- COMET LEAP methodology

****Theoretical Impact:****

- Flourishing measurement
- Evidence-based interventions
- Wisdom preservation
- Bayesian modeling
- 30-day iteration

Capability Maturity Matrix

Capability	Theoretical	Prototype	Proven	Production
Local-First AI Control Plane	■	■	■	■
Agent Orchestration	■	■	■	■
Game Development Tools	■	■	■	■
Business Services	■	■	■	■
Symbolic Sandbox	■	■	■	■
ASI Arc Reactor MK1	■	■	■	■
COMET LEAP Framework	■	■	■	■
Windsurf Hooks	■	■	■	■
Convergence Loop	■	■	■	■
Human Flourishing Frameworks	■	■	■	■

Integration Maturity Matrix

Integration	Theoretical	Prototype	Proven	Production
Lantern OS + COMET LEAP	■	■	■	■
Agent Orchestration + Symbolic Sandbox	■	■	■	■



Integration	Theoretical	Prototype	Proven	Production
Business Services + Payment Bridge	■	■	■	■
Game Development + Cloud Access	■	■	■	■
Safety Hooks + Audit Logging	■	■	■	■

Next Steps

Immediate (30 Days)

1. ****ASI Arc Reactor MK1****: Execute \$1000 founding seat demos for evidence
2. ****MCP Canary****: Validate MCP canary for automation safety
3. ****Convergence Loop****: Run convergence loop and fix first 4 issues
4. ****Business Services****: Complete secrets/env review for business services
5. ****Symbolic Sandbox****: Complete rights review for content dependencies

Short-term (90 Days)

1. ****COMET LEAP Integration****: Integrate COMET LEAP framework into Lantern OS
2. ****Agent Orchestration****: Complete agent orchestration with safety hooks
3. ****Business Automation****: Complete business service automation
4. ****Human Flourishing Metrics****: Develop flourishing metrics
5. ****Evidence-Based Promotion****: Implement evidence-based promotion system

Long-term (1 Year)

1. ****Production Readiness****: Move core capabilities to production
2. ****Human Trial Demo****: Execute human trial demo with evidence
3. ****Founder Wisdom****: Extract and use founder wisdom in decisions
4. ****Flourishing Frameworks****: Deploy human flourishing frameworks
5. ****ASI Pattern Integration****: Complete ASI pattern integration with evidence

Summary

The alex-place and human-flourishing-frameworks organizations have developed comprehensive theoretical and novel capabilities across local-first AI control, agent orchestration, game development, business services, and human flourishing frameworks.

****Key Innovations:****

- ASI Arc Reactor MK1 with strict claim boundaries
- COMET LEAP framework for founder wisdom
- Windsurf hooks for safety validation



- Convergence loop for evidence-based decisions
- Human flourishing frameworks

****Current Maturity:****

- Theoretical: All capabilities
- Prototype: Most capabilities
- Proven: Core capabilities (game dev, convergence loop, safety hooks)
- Production: Limited (game development, convergence loop)

****Next Focus:****

- Evidence collection for ASI Arc Reactor MK1
- MCP canary validation
- Business services secrets/env review
- COMET LEAP integration
- Human flourishing metrics development

****Report Generated:**** 2026-05-30

****Organizations:**** alex-place, human-flourishing-frameworks

****Scanner:**** Cascade (Lantern OS Agent)

****Status:**** Comprehensive capability mapping complete

