

MCP Canary Test Report

Generated: 2026-05-31T03:31:29
Engine: MCP Canary Test Engine v1.0
Status: ■ SAFE TO PROCEED

SIMPLE ANSWER

MCP Canary validation complete. **25** of 29 sources verified. **1** held. **0** warnings. All sources safe - no tool execution enabled.

WHAT IT ACTUALLY DOES

The MCP Canary Test Engine validates all Model Context Protocol sources before any automation or tool execution:

- 1. **Tests connectivity** - Verifies each source can be reached
- 2. **Checks safety flags** - Confirms ■llowToolExecution is disabled
- 3. **Validates local-only** - Ensures local sources stay local
- 4. **Verifies auth** - Confirms auth requirements are documented
- 5. **Blocks on warnings** - Prevents automation if any source has tool execution enabled

This is a **safety gate** - not a capability test. The engine asks "is it safe?" not "does it work?"

EVIDENCE / SOURCE DISCIPLINE

MCP Sources Manifest: D:\tmp\lantern-os\manifests\lantern-mcp-sources.json
Total Sources: 29

By Transport:

- Source repos: 3
- Local filesystem: 2
- Web scrape: 18
- REST API: 3
- HTTP JSON-RPC: 3

Safety Verification:

- All sources checked for ■llowToolExecution
- Local sources verified for localOnly flag
- Auth requirements documented
- Rate limits respected for public APIs

PROVEN / HELD / LOCAL-ONLY

Verified Sources (Safe to Reference):

- ■ **GM Agent Orchestrator (GitHub)** - read-only-verification
- ■ **U.S. Census BTOS** - reference-only
- ■ **U.S. Chamber Small Business AI Report 2025** - reference-only
- ■ **Pew Research Homeschooling Data** - reference-only
- ■ **Johns Hopkins Homeschool Growth 2024-2025** - reference-only
- ■ **EdChoice K-12 Tutoring Survey** - reference-only
- ■ **Business.com SMB AI Outlook 2026** - reference-only
- ■ **Kalshi Public Markets API** - public-api-verified
- ■ **Kalshi API Documentation** - reference-only
- ■ **Ko-fi Support Page** - reference-only
- ■ **Human Flourishing Frameworks Scan (GitHub)** - read-only-verification
- ■ **HFF Scan Local Clone** - local-verified

- ■ **Lantern OS Repository** - read-only-verification
- ■ **Lantern OS Local Workspace** - local-verified
- ■ **Kickstarter** - reference-only
- ■ **Patreon** - reference-only
- ■ **GoFundMe** - reference-only
- ■ **Gumroad** - reference-only
- ■ **Etsy** - reference-only
- ■ **Upwork** - reference-only
- ■ **GitHub Sponsors** - reference-only
- ■ **Itch.io** - reference-only
- ■ **Steamworks (Steam Direct)** - reference-only
- ■ **GOG Developer Portal** - reference-only
- ■ **Internet Archive** - public-api-verified

Held Sources (Require Action):

- ■■ **GM Agent Orchestrator Local** - skipped
- ■■ **AWS CLI MCP (Local)** - skipped
- ■■ **AWS Bedrock MCP (Local)** - skipped
- ■■ **GitHub API** - auth-required

Local-Only:

- This report is local evidence only
- CanProceed flag is advisory
- Actual tool execution requires operator approval
- No automated execution without explicit sign-off

NEXT SAFE ACTION

MCP sources validated - safe to reference

Maintenance:

- 1. Re-run canary after adding new MCP sources
- 2. Verify held sources weekly
- 3. Check for new transport types
- 4. Update safety boundaries if architecture changes

VALIDATION PATH

- [] Review canary test results
- [] Verify ■■SourcesSafe: true
- [] Check held sources are intentional
- [] Confirm no tool execution without approval
- [] Operator sign-off for MCP-based automation
- [] Archive this report in manifests/validation/

Generated by MCP Canary Test Engine

Skill Reference: skills/asi-arc-reactor-mk1/SKILL.md

Style Reference: docs/ORION-MOOKMANREPORT4-STYLE.md